Chujie Zheng 郑楚杰

Contact Information

Address(es)	Room 4-504, FIT Building, Tsinghua University, Beijing 100084, China
Email	chujiezhengchn@gmail.com
Homepage	https://chujiezheng.github.io/
Google Scholar	https://scholar.google.com/citations?user=55zBNgUAAAAJ

Education

- **Ph.D. student**, Department of Computer Science and Technology, Tsinghua University   Sep 2020 – Present
  Advisor: Minlie Huang
- **Visiting Scholar**, Computer Science Department, UCLA   Nov 2023 – Jun 2024
  Advisor: Nanyun (Violet) Peng
- **B.S.** in Physics, Tsinghua University   Aug 2016 – Jun 2020

Research Overview and Highlights

- I have a broad research interest in **building trustworthy AI systems**, with the current focus on **LLM alignment** (preprint, ICLR 2024 Spotlight, ICML 2024). My ultimate research goal is to **advance and oversee AI systems in a scalable way (with minimal human intervention) and ensure they operate transparently and responsibly**.
- Previously, I conducted research on **LLMs for social good**, covering the topics of emotional support dialogue systems (ACL 2021, ACL 2023) and empathetic dialogue generation (ACL 2021, AAAI 2022, ACL 2023).
- In my early research career, I built a series of **popular NLP datasets**, including ChID (ACL 2019), KDConv (ACL 2020), ESConv (ACL 2021), CDConv (EMNLP 2022), DiaSafety (ACL 2022), and COLD (EMNLP 2022).
- Google Scholar citations **1050+**, h-index **16**, i10-index **18**.

Selected Projects/Papers

1. **Weak-to-Strong Extrapolation Expedites Alignment**
   Chujie Zheng, Ziqi Wang, Heng Ji, Minlie Huang, Nanyun Peng
   arXiv:2404.16792 (10K+ downloads on 🤗HuggingFace in 2 weeks) (12 citations)
   [paper] [repo] [🤗 HuggingFace]

2. **On Prompt-Driven Safeguarding for Large Language Models**
   Chujie Zheng, Fan Yin, Hao Zhou, Fandong Meng, Jie Zhou, Kai-Wei Chang, Minlie Huang, Nanyun Peng
   ICML 2024 || Secure and Trustworthy LLM Workshop @ ICLR 2024 (Oral: 5%)
   [paper] [repo]

3. **Chat Templates for 🤗 HuggingFace Large Language Models**
   Chujie Zheng
   GitHub Repository (310+ stars)

[repo]

4. **Large Language Models Are Not Robust Multiple Choice Selectors**
   Chujie Zheng, Hao Zhou, Fandong Meng, Jie Zhou, Minlie Huang
   ICLR 2024 (Spotlight: 5%) (68 citations)
   [paper] [repo]

5. **Click: Controllable Text Generation with Sequence Likelihood Contrastive Learning**
   Chujie Zheng, Pei Ke, Zheng Zhang, Minlie Huang
   ACL 2023 Findings
   [paper] [repo]

6. **Towards Emotional Support Dialog Systems**
   Siyang Liu*, **Chujie Zheng***, Orianna Demasi, Sahand Sabour, Yu Li, Zhou Yu, Yong Jiang, Minlie Huang (*: Equal contribution)
   ACL 2021 (Oral) (182 citations)
   [paper] [repo]

Other Papers

7. **AugESC: Dialogue Augmentation with Large Language Models for Emotional Support Conversation**
   Chujie Zheng, Sahand Sabour, Jiaxin Wen, Zheng Zhang, Minlie Huang
   ACL 2023 Findings

8. **CASE: Aligning Coarse-to-Fine Cognition and Affection for Empathetic Response Generation**
   Jinfeng Zhou*, **Chujie Zheng***, Bo Wang, Zheng Zhang, Minlie Huang
   ACL 2023 (Oral)

9. **EVA2.0: Investigating Open-domain Chinese Dialogue Systems with Large-scale Pre-training**
   Yuxian Gu*, Jiaxin Wen*, Hao Sun*, Yi Song, Pei Ke, **Chujie Zheng**, Zheng Zhang, Jianzhu Yao, Lei Liu, Xiaoyan Zhu, Minlie Huang
   Machine Intelligence Research 2023

10. **CDConv: A Benchmark for Contradiction Detection in Chinese Conversations**
    Chujie Zheng*, Jinfeng Zhou*, Yinhe Zheng, Libiao Peng, Zhen Guo, Wenquan Wu, Zhengyu Niu, Hua Wu, Minlie Huang
    EMNLP 2022 (Oral)

11. **COLD: A Benchmark for Chinese Offensive Language Detection**
    Jiawen Deng*, Jingyan Zhou*, Hao Sun, **Chujie Zheng**, Fei Mi, Helen Meng, Minlie Huang
    EMNLP 2022 (Oral)

12. **On the Safety of Conversational Models: Taxonomy, Dataset, and Benchmark**
    Hao Sun*, Guangxuan Xu*, Jiawen Deng, Jiale Cheng, **Chujie Zheng**, Hao Zhou, Nanyun Peng, Xiaoyan Zhu, Minlie Huang
    ACL 2022 Findings

13. **CEM: Commonsense-aware Empathetic Response Generation**
    Sahand Sabour, **Chujie Zheng**, Minlie Huang
    AAAI 2022 (Oral)

14. **Exploring Prompt-based Few-shot Learning for Grounded Dialog Generation**
    Chujie Zheng, Minlie Huang
    arXiv:2109.06513

15. **EVA: An Open-Domain Chinese Dialogue System with Large-Scale Generative Pre-Training**
    Hao Zhou*, Pei Ke*, Zheng Zhang*, Yuxian Gu, Yinhe Zheng, **Chujie Zheng**, Yida Wang, Chen Henry Wu, Hao Sun, Xiaocong Yang, Bosi Wen, Xiaoyan Zhu, Minlie Huang, Jie Tang
    arXiv:2108.01547

16. **CoMAE: A Multi-factor Hierarchical Framework for Empathetic Response Generation**
    Chujie Zheng, Yong Liu, Wei Chen, Yongcai Leng, Minlie Huang
    ACL 2021 Findings

17. **PsyQA: A Chinese Dataset for Generating Long Counseling Text for Mental Health Support**
    Hao Sun*, Zhenru Lin*, **Chujie Zheng**, Siyang Liu, Minlie Huang
    ACL 2021 Findings

18. **Difference-aware Knowledge Selection for Knowledge-grounded Conversation Generation**
    Chujie Zheng, Yunbo Cao, Daxin Jiang, Minlie Huang
    EMNLP 2020 Findings

19. **KdConv: A Chinese Multi-domain Dialogue Dataset Towards Multi-turn Knowledge-driven Conversation**
    Hao Zhou*, **Chujie Zheng***, Kaili Huang, Minlie Huang, Xiaoyan Zhu
    ACL 2020

20. **ChID: A Large-scale Chinese IDiom Dataset for Cloze Test**
    Chujie Zheng, Minlie Huang, Aixin Sun
    ACL 2019

Selected Awards and Honors

- Schlumberger Scholarship, Tsinghua University   2023
- Comprehensive Merit Scholarship, Tsinghua University   2022
- Outstanding Undergraduate, Tsinghua University   2020
- China National Scholarship (Top 2/100)   2019
- Comprehensive Merit Scholarship, Tsinghua University   2018

Academic Services

- **Area Chair:** ACL (24), EMNLP (24), ACL Rolling Review (24)
- **Reviewer:** NeurIPS (24), ICML (24), ACL (22/23), EMNLP (21/22), NAACL (24), EACL (23), ACL Rolling Review (21/22/23), CogSci (24), AAAI (22/23)